

```

In [24]: #cell 1
import os
import random
import numpy as np
import pandas as pd
import tensorflow as tf
from tensorflow.keras import layers, models
from tensorflow.keras.applications import ResNet50
from tensorflow.keras.optimizers import Adam
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns
import time

# --- Configuration ---
CSV_PATH = '/cluster/home/miolate21/FER_biasmitigation1/results/inferred_com
IMAGE_SHAPE = (224, 224, 1)
INPUT_SHAPE_MODEL = (224, 224, 3)
BATCH_SIZE = 32
SEED = 42
AUTOTUNE = tf.data.AUTOTUNE
NUM_CLASSES = 7
INITIAL_EPOCHS = 5
FINE_TUNE_EPOCHS = 15 # ajustar si hay que
TOTAL_EPOCHS = INITIAL_EPOCHS + FINE_TUNE_EPOCHS

# Learning Rates
INITIAL_LR = 1e-4
FINE_TUNE_LR = 1e-5

# Model Saving
MODEL_SAVE_DIR = "/cluster/home/miolate21/FER_biasmitigation1/models"

MODEL_SAVE_PATH_FEATURE = os.path.join(MODEL_SAVE_DIR, "resnet50_fer_feature
MODEL_SAVE_PATH_FINETUNE = os.path.join(MODEL_SAVE_DIR, "resnet50_fer_finetu
# Ensure model directory exists
os.makedirs(MODEL_SAVE_DIR, exist_ok=True)

# --- Reproducibility ---
os.environ['PYTHONHASHSEED'] = str(SEED)
random.seed(SEED)
np.random.seed(SEED)
tf.random.set_seed(SEED)
# tf.keras.utils.set_random_seed(SEED) # Use this for newer TF versions if r

print(f"Using SEED: {SEED}")
print(f"Input Shape expected by NPY loader: {IMAGE_SHAPE}")
print(f"Model Input Shape (after processing): {INPUT_SHAPE_MODEL}")
print(f"Number of Classes: {NUM_CLASSES}")
print(f"Initial LR: {INITIAL_LR}, Fine-tune LR: {FINE_TUNE_LR}")
print(f"Initial Epochs: {INITIAL_EPOCHS}, Fine-tune Epochs: {FINE_TUNE_EPOCHS}")
print(f"Model save path: {MODEL_SAVE_PATH}")

```

Using SEED: 42
 Input Shape expected by NPY loader: (224, 224, 1)
 Model Input Shape (after processing): (224, 224, 3)
 Number of Classes: 7
 Initial LR: 0.0001, Fine-tune LR: 1e-05
 Initial Epochs: 5, Fine-tune Epochs: 15
 Model save path: /cluster/home/miolate21/FER_biasmitigation1/models/resnet50_baseline_224.keras

In [25]: `#cell 2`

```
print(f"Loading CSV from: '{CSV_PATH}'")
df = None
idx_to_emotion = None
CLASS_NAMES = None

try:
    if not os.path.exists(CSV_PATH):
        raise FileNotFoundError(f"CSV file not found at {CSV_PATH}")

    df = pd.read_csv(CSV_PATH)

    if 'Unnamed: 0' in df.columns:
        df = df.drop(columns=['Unnamed: 0'])

    print(f"Loaded {len(df)} records.")

    required_columns = ['image_path', 'emotion']
    if not all(col in df.columns for col in required_columns):
        raise ValueError(f"CSV must contain columns: {required_columns}")

    # Create label mappings
    emotion_labels = sorted(df['emotion'].unique())
    if len(emotion_labels) != NUM_CLASSES:
        print(f"Warning: Found {len(emotion_labels)} unique emotions, but ex

    emotion_to_idx = {label: idx for idx, label in enumerate(emotion_labels)}
    idx_to_emotion = {idx: label for label, idx in emotion_to_idx.items()}
    CLASS_NAMES = list(idx_to_emotion.values())

    df['label'] = df['emotion'].map(emotion_to_idx)

    if df['label'].isnull().any():
        print("Warning: Some emotions could not be mapped to labels.")
        print(df[df['label'].isnull()]['emotion'].unique())

    print("\nLabel map created:")
    print(f"Emotion to Index: {emotion_to_idx}")
    print(f"Index to Emotion: {idx_to_emotion}")
    print(f"Class Names: {CLASS_NAMES}")
    print(f"\nDataFrame head:\n{df.head()}")

except FileNotFoundError as e:
    print(f"ERROR: {e}")
except ValueError as e:
```

```

    print(f"ERROR: {e}")
except Exception as e:
    print(f"An error occurred loading or processing the CSV: {e}")
    df = None # Ensure df is None on error

# Critical check before proceeding
if df is None or df.empty:
    print("\nCritical Error: DataFrame is empty or None. Cannot proceed.")
    # Stop execution or handle appropriately
    assert False, "DataFrame loading failed."
elif df['label'].isnull().any():
    print("\nCritical Error: DataFrame contains rows with null labels after")
    assert False, "Label mapping failed for some rows."

```

Loading CSV from: '/cluster/home/miolate21/FER_biasmitigation1/results/infer
red_combined_224.csv'

Loaded 41476 records.

Label map created:

Emotion to Index: {'angry': 0, 'disgust': 1, 'fear': 2, 'happy': 3, 'neutra
l': 4, 'sad': 5, 'surprise': 6}

Index to Emotion: {0: 'angry', 1: 'disgust', 2: 'fear', 3: 'happy', 4: 'neut
ral', 5: 'sad', 6: 'surprise'}

Class Names: ['angry', 'disgust', 'fear', 'happy', 'neutral', 'sad', 'surpri
se']

DataFrame head:

	image_path	dataset	emotion	label
0	/cluster/home/miolate21/FER_biasmitigation1/da...	FER2013	neutral	4
1	/cluster/home/miolate21/FER_biasmitigation1/da...	FER2013	neutral	4
2	/cluster/home/miolate21/FER_biasmitigation1/da...	FER2013	neutral	4
3	/cluster/home/miolate21/FER_biasmitigation1/da...	FER2013	neutral	4
4	/cluster/home/miolate21/FER_biasmitigation1/da...	FER2013	neutral	4

In [8]: #cell 3

```

if df is not None and not df.empty:
    path_column_to_check = 'image_path'
    print(f"\nChecking for .npz files listed in '{path_column_to_check}'...")

    start_time = time.time()
    exists = df[path_column_to_check].apply(os.path.exists)
    end_time = time.time()

    num_missing = len(df) - exists.sum()
    if num_missing > 0:
        print(f"\nFound {num_missing} missing .npz files.")
        missing_files_df = df[~exists]
        print("Missing file examples (first 5):")
        print(missing_files_df[path_column_to_check].head().tolist())

```

```

else:
    print("All listed .npy files are present.")
else:
    print("No data loaded skipping file check.")

```

Checking for .npy files listed in 'image_path'...
All listed .npy files are present.

In [11]: #cell 4 Train/Validation Split

```

train_df = None
val_df = None

if df is not None and not df.empty:
    if 'label' not in df.columns or df['label'].isnull().any():
        print("ERROR: 'label' column is missing")
    else:
        print("\nSplitting dataset into training and validation sets...")

        try:
            train_df, val_df = train_test_split(
                df,
                test_size=0.2,          # 20% goes to validation
                stratify=df['label'],
                random_state=SEED
            )
            print(f"Training set size: {len(train_df)} samples")
            print(f"Validation set size: {len(val_df)} samples")

            print("\nTraining set label distribution (%):")
            print((train_df['label'].value_counts(normalize=True).sort_index()

            print("\nValidation set label distribution (%):")
            print((val_df['label'].value_counts(normalize=True).sort_index()

        except Exception as e:
            print(f"Split failed: {e}")
            print("Check your label distribution.")
            train_df, val_df = pd.DataFrame(), pd.DataFrame()
    else:
        print("DataFrame is empty or None. Skipping split step.")

```

Splitting dataset into training and validation sets...

Training set size: 33180 samples

Validation set size: 8296 samples

Training set label distribution (%):

label

0 11.42%

1 2.89%

2 10.61%

3 29.03%

4 18.67%

5 16.47%

6 10.91%

Name: proportion, dtype: object

Validation set label distribution (%):

label

0 11.42%

1 2.89%

2 10.61%

3 29.03%

4 18.67%

5 16.48%

6 10.91%

Name: proportion, dtype: object

In [26]: *# Cell 5 Dataset Loader and Builder*

--- Data Augmentation Layer ---

```
data_augmentation = tf.keras.Sequential([
    layers.RandomFlip("horizontal", seed=SEED),
    layers.RandomRotation(0.1, seed=SEED),
    layers.RandomZoom(0.1, seed=SEED),
], name="data_augmentation")
```

```
def _load_and_process_npy(path_bytes, label_int):
```

```
    path_str = path_bytes.numpy().decode('utf-8')
```

```
    try:
```

```
        img = np.load(path_str)
```

--- Data Type and Range ---

```
    img = img.astype(np.float32)
```

```
    if img.max() <= 1.0: img = img * 255.0
```

```
    if img.ndim == 2:
```

```
        img = np.expand_dims(img, axis=-1) # Add channel: (H, W, 1)
```

```
    if img.shape != IMAGE_SHAPE:
```

```
        if img.shape == IMAGE_SHAPE[:-1]:
```

```
            img = np.expand_dims(img, axis=-1)
```

```
        else:
```

```
            raise ValueError(f"Unexpected shape {img.shape} for image {p
```

```

        return img, np.int32(label_int)

    except FileNotFoundError:
        print(f"ERROR (_load_and_process_npy): File not found: {path_str}. F
        return np.zeros(IMAGE_SHAPE, dtype=np.float32), np.int32(label_int)
    except ValueError as ve: # Catch specific shape errors
        print(f"ERROR (_load_and_process_npy): Shape error for {path_str}:
        return np.zeros(IMAGE_SHAPE, dtype=np.float32), np.int32(label_int)
    except Exception as e:
        print(f"ERROR (_load_and_process_npy): Problem loading {path_str}: {
        return np.zeros(IMAGE_SHAPE, dtype=np.float32), np.int32(label_int)

@tf.function
def load_npy_tf_wrapper(path_tensor, label_tensor):
    image, label = tf.py_function(
        _load_and_process_npy,
        inp=[path_tensor, label_tensor],
        Tout=(tf.float32, tf.int32)
    )
    image.set_shape(IMAGE_SHAPE)
    label.set_shape(())
    return image, label

# --- Updated Dataset Builder ---
def build_tf_dataset(df_subset, shuffle=True, augment=False, batch_size=BATCH_SIZE):
    if df_subset is None or df_subset.empty:
        print("DataFrame subset is empty, returning empty dataset.")
        return tf.data.Dataset.from_tensor_slices(
            (tf.constant([], dtype=tf.string), tf.constant([], dtype=tf.int32))
        ).map(load_npy_tf_wrapper).batch(batch_size).prefetch(buffer_size=AUTOTUNE)

    paths = df_subset['image_path'].values
    labels = df_subset['label'].values.astype(np.int32)

    ds = tf.data.Dataset.from_tensor_slices((paths, labels))

    if shuffle:
        ds = ds.shuffle(buffer_size=len(df_subset), seed=SEED, reshuffle_each_batch=True)

    ds = ds.map(load_npy_tf_wrapper, num_parallel_calls=AUTOTUNE)

    # --- Apply Augmentation ---
    if augment:
        ds = ds.map(lambda x, y: (data_augmentation(x, training=True), y),
                    num_parallel_calls=AUTOTUNE)

    # Batch dataset
    ds = ds.batch(batch_size, drop_remainder=False)

    # Prefetch for performance
    ds = ds.prefetch(buffer_size=AUTOTUNE)

    return ds

# --- Build the Datasets ---
print("Building training dataset...")

```

```

train_ds = build_tf_dataset(train_df, shuffle=True, augment=True, batch_size=
print("\nBuilding validation dataset...")
val_ds = build_tf_dataset(val_df, shuffle=False, augment=False, batch_size=E
print("\nDatasets ready. Train/Val tf.data.Dataset objects created.")

```

Building training dataset...

Building validation dataset...

Datasets ready. Train/Val tf.data.Dataset objects created.

```

In [46]: #cell 6 Model Building

def build_resnet50_transfer(input_shape_gray=IMAGE_SHAPE, num_classes=NUM_CLASSES):

    # --- Input Layer ---
    inputs = layers.Input(shape=input_shape_gray, name="input_layer")

    # --- Preprocessing ---
    x = layers.Concatenate(axis=-1, name="grayscale_to_rgb")([inputs, inputs])
    x = tf.keras.applications.resnet50.preprocess_input(x)

    # --- Base ResNet50 Model ---
    base_model = ResNet50(
        include_top=False,
        weights='imagenet',
        input_shape=INPUT_SHAPE_MODEL
    )
    # freeze base model
    base_model.trainable = False
    print(f"Base ResNet50 model loaded. Initial trainable status: {base_model.trainable}")

    # --- Connect Preprocessing to Base Model ---
    x = base_model(x, training=False)

    # --- Classifier Head ---
    x_head = layers.GlobalAveragePooling2D(name="global_avg_pool")(x)
    x_head = layers.BatchNormalization()(x_head)
    x_head = layers.Dense(128, activation='relu', name="dense_head")(x_head)
    x_head = layers.Dropout(0.5, name="dropout_head")(x_head)
    outputs = layers.Dense(num_classes, activation='softmax', name="output_logits")

    # --- Create the final model ---
    model = models.Model(inputs=inputs, outputs=outputs, name="ResNet50_Transfer")

    return model

# --- Instantiate the model ---
print("Building ResNet50 transfer learning model...")
model = build_resnet50_transfer(input_shape_gray=IMAGE_SHAPE, num_classes=NUM_CLASSES)

print("\nInitial Model Summary (Base Frozen):")
model.summary(line_length=120)

```

```
# Optional: Verify trainable/non-trainable counts programmatically
total_params = model.count_params()
trainable_count = sum([tf.keras.backend.count_params(w) for w in model.trainable_weights])
non_trainable_count = sum([tf.keras.backend.count_params(w) for w in model.non_trainable_weights])
print(f"\nTotal parameters: {total_params:,}")
print(f"Trainable parameters (initially): {trainable_count:,}")
print(f"Non-trainable parameters (initially): {non_trainable_count:,}")
```

Building ResNet50 transfer learning model...

Base ResNet50 model loaded. Initial trainable status: False

Initial Model Summary (Base Frozen):

Model: "ResNet50_Transfer_FER"

Layer (type)	Output Shape
input_layer (InputLayer)	(None, 224, 224, 1)
grayscale_to_rgb (Concatenate)	(None, 224, 224, 3)
get_item_30 (GetItem)	(None, 224, 224)
get_item_31 (GetItem)	(None, 224, 224)
get_item_32 (GetItem)	(None, 224, 224)
stack_10 (Stack)	(None, 224, 224, 3)
add_10 (Add)	(None, 224, 224, 3)
resnet50 (Functional)	(None, 7, 7, 2048)
global_avg_pool (GlobalAveragePooling2D)	(None, 2048)
batch_normalization_9 (BatchNormalization)	(None, 2048)
dense_head (Dense)	(None, 128)
dropout_head (Dropout)	(None, 128)
output_layer (Dense)	(None, 7)

Total params: 23,859,079 (91.02 MB)

Trainable params: 267,271 (1.02 MB)

Non-trainable params: 23,591,808 (90.00 MB)

Total parameters: 23,859,079

Trainable parameters (initially): 267,271

Non-trainable parameters (initially): 23,591,808

In [47]: *#cell 7 initial compilation*

```
print("Compiling the model for initial feature extraction...")
model.compile(
    optimizer=Adam(learning_rate=INITIAL_LR),
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy']
)
print(f"Model compiled successfully with LR={INITIAL_LR}.")
```

Compiling the model for initial feature extraction...

Model compiled successfully with LR=0.0001.

In [39]: *#cell 8 Initial Feature Extraction Training*

```
print(f"\n--- Starting Initial Training (Feature Extraction) for {INITIAL_EPOCHS} epochs")
print(f"Saving best model checkpoints during this phase to: {MODEL_SAVE_PATH}")

callbacks_initial = [
    tf.keras.callbacks.ModelCheckpoint(
        filepath=MODEL_SAVE_PATH,
        save_best_only=True,
        monitor='val_accuracy',
        mode='max',
        verbose=1
    ),
    tf.keras.callbacks.EarlyStopping(
        monitor='val_accuracy',
        patience=3,
        verbose=1,
        restore_best_weights=True
    )
]

if 'train_ds' not in globals() or 'val_ds' not in globals() or train_ds is None or val_ds is None:
    raise NameError("ERROR: 'train_ds' or 'val_ds' not defined or is None.")











if tf.data.experimental.cardinality(train_ds).numpy() == 0 or \
    tf.data.experimental.cardinality(val_ds).numpy() == 0:
    raise ValueError("ERROR: Training or validation dataset is empty.")

history_initial = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=INITIAL_EPOCHS,
    callbacks=callbacks_initial,
    verbose=1
)

print("\n--- Initial model training (feature extraction phase) finished. ---")

# Keep track of the best validation accuracy achieved in this phase
initial_best_val_acc = max(history_initial.history['val_accuracy'])
print(f"Best validation accuracy during initial phase: {initial_best_val_acc}")
```

```

--- Starting Initial Training (Feature Extraction) for 5 epochs ---
Saving best model checkpoints during this phase to: /cluster/home/miolate21/
FER_biasmitigation1/models/resnet50_baseline_224.keras
Epoch 1/5
1037/1037  0s 62ms/step - accuracy: 0.2662 - loss: 2.232
8
Epoch 1: val_accuracy improved from -inf to 0.39947, saving model to /cluster/home/miolate21/FER_biasmitigation1/models/resnet50_baseline_224.keras
1037/1037  89s 81ms/step - accuracy: 0.2663 - loss: 2.23
26 - val_accuracy: 0.3995 - val_loss: 1.6045
Epoch 2/5
1035/1037  0s 60ms/step - accuracy: 0.3640 - loss: 1.747
7
Epoch 2: val_accuracy improved from 0.39947 to 0.43599, saving model to /cluster/home/miolate21/FER_biasmitigation1/models/resnet50_baseline_224.keras
1037/1037  78s 75ms/step - accuracy: 0.3640 - loss: 1.74
76 - val_accuracy: 0.4360 - val_loss: 1.5334
Epoch 3/5
1037/1037  0s 60ms/step - accuracy: 0.3928 - loss: 1.610
7
Epoch 3: val_accuracy improved from 0.43599 to 0.46601, saving model to /cluster/home/miolate21/FER_biasmitigation1/models/resnet50_baseline_224.keras
1037/1037  77s 75ms/step - accuracy: 0.3928 - loss: 1.61
07 - val_accuracy: 0.4660 - val_loss: 1.4338
Epoch 4/5
1037/1037  0s 59ms/step - accuracy: 0.4226 - loss: 1.538
6
Epoch 4: val_accuracy did not improve from 0.46601
1037/1037  76s 73ms/step - accuracy: 0.4226 - loss: 1.53
86 - val_accuracy: 0.4594 - val_loss: 1.4307
Epoch 5/5
1036/1037  0s 60ms/step - accuracy: 0.4320 - loss: 1.493
9
Epoch 5: val_accuracy improved from 0.46601 to 0.47191, saving model to /cluster/home/miolate21/FER_biasmitigation1/models/resnet50_baseline_224.keras
1037/1037  77s 74ms/step - accuracy: 0.4320 - loss: 1.49
39 - val_accuracy: 0.4719 - val_loss: 1.3871
Restoring model weights from the end of the best epoch: 5.

--- Initial model training (feature extraction phase) finished. ---
Best validation accuracy during initial phase: 0.4719

```

In [48]: *# Cell 9 Fine-tuning Setup*

```

print("\n--- Preparing for Fine-Tuning ---")

# --- Unfreeze Layers ---
try:
    base_model = model.get_layer('resnet50')
except ValueError:
    print("ERROR: Could not find layer named 'resnet50'.")
    possible_names = [l.name for l in model.layers if 'resnet50' in l.name.lower()]
    if not possible_names:
        raise ValueError("Could not find the ResNet50 base layer in the model")
    base_model_name = possible_names[0]
    print(f"Found base layer with name: {base_model_name}")

```

```

base_model = model.get_layer(base_model_name)

base_model.trainable = True
print(f"Base ResNet50 model trainable status set to: {base_model.trainable}")

fine_tune_at_layer_name = 'conv5_block1_out'

try:
    fine_tune_from_index = [i for i, layer in enumerate(base_model.layers) if layer.name == fine_tune_at_layer_name]
    print(f"Found fine-tune layer '{fine_tune_at_layer_name}' at index {fine_tune_from_index}")

    # Freeze all layers before the fine_tune_from_index
    for layer in base_model.layers[:fine_tune_from_index]:
        layer.trainable = False
    print(f"Froze layers in base model up to index {fine_tune_from_index}.")

except IndexError:
    print(f"Warning: Could not find layer named '{fine_tune_at_layer_name}'")
    print("Available layer names in base_model:", [layer.name for layer in base_model.layers])
    print("Proceeding with the entire base model potentially trainable .")

# --- Re-compile the Model ---
print(f"\nRe-compiling model for fine-tuning with LR={FINE_TUNE_LR}...")
model.compile(
    optimizer=Adam(learning_rate=FINE_TUNE_LR),
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy']
)
print("Model re-compiled successfully.")

# --- Verify Trainable Parameters ---
print("\nModel Summary (During Fine-Tuning):")
model.summary(line_length=120)

trainable_count_ft = sum([tf.keras.backend.count_params(w) for w in model.trainable_weights])
non_trainable_count_ft = sum([tf.keras.backend.count_params(w) for w in model.non_trainable_weights])
print(f"\nTotal parameters: {model.count_params():,}")
print(f"Trainable parameters (fine-tuning): {trainable_count_ft:,}")
print(f"Non-trainable parameters (fine-tuning): {non_trainable_count_ft:,}")

# Check difference
print(f"Number of parameters unfrozen in base model: {trainable_count_ft - non_trainable_count_ft}")

```

--- Preparing for Fine-Tuning ---

Base ResNet50 model trainable status set to: True

Found fine-tune layer 'conv5_block1_out' at index 154 in the base model.

Froze layers in base model up to index 154.

Re-compiling model for fine-tuning with LR=1e-05...

Model re-compiled successfully.

Model Summary (During Fine-Tuning):

Model: "ResNet50_Transfer_FER"

Layer (type)	Output Shape
input_layer (InputLayer)	(None, 224, 224, 1)
grayscale_to_rgb (Concatenate)	(None, 224, 224, 3)
get_item_30 (GetItem)	(None, 224, 224)
get_item_31 (GetItem)	(None, 224, 224)
get_item_32 (GetItem)	(None, 224, 224)
stack_10 (Stack)	(None, 224, 224, 3)
add_10 (Add)	(None, 224, 224, 3)
resnet50 (Functional)	(None, 7, 7, 2048)
global_avg_pool (GlobalAveragePooling2D)	(None, 2048)
batch_normalization_9 (BatchNormalization)	(None, 2048)
dense_head (Dense)	(None, 128)
dropout_head (Dropout)	(None, 128)
output_layer (Dense)	(None, 7)

Total params: 23,859,079 (91.02 MB)

Trainable params: 9,198,599 (35.09 MB)

Non-trainable params: 14,660,480 (55.93 MB)

Total parameters: 23,859,079

Trainable parameters (fine-tuning): 9,198,599

Non-trainable parameters (fine-tuning): 14,660,480

Number of parameters unfrozen in base model: 9,194,503

In [49]: *#Cell 10 Fine-Tuning Training*

```

print(f"\n--- Starting Fine-Tuning Training for up to {FINE_TUNE_EPOCHS} epochs")
print(f"Continuing training from epoch {INITIAL_EPOCHS}.")
print(f"Saving best model checkpoints during this phase to: {MODEL_SAVE_PATH}")

callbacks_finetune = [
    tf.keras.callbacks.ModelCheckpoint(
        filepath=MODEL_SAVE_PATH,
        save_best_only=True,
        monitor='val_accuracy',
        mode='max',
        verbose=1
    ),

```

```

tf.keras.callbacks.EarlyStopping(
    monitor='val_accuracy',
    patience=5,
    verbose=1,
    restore_best_weights=True
)

]

if 'train_ds' not in globals() or 'val_ds' not in globals() or train_ds is None:
    raise NameError("ERROR: 'train_ds' or 'val_ds' not defined or is None")
if tf.data.experimental.cardinality(train_ds).numpy() == 0 or \
    tf.data.experimental.cardinality(val_ds).numpy() == 0:
    raise ValueError("ERROR: Training or validation dataset is empty for fine-tuning")

TOTAL_EPOCHS = INITIAL_EPOCHS + FINE_TUNE_EPOCHS
print(f"Training will run until epoch {TOTAL_EPOCHS} unless stopped early.")

history_fine_tune = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=TOTAL_EPOCHS,
    initial_epoch=INITIAL_EPOCHS,
    callbacks=callbacks_finetune,
    verbose=1
)
















print("\n--- Fine-tuning training phase finished. ---")

if history_fine_tune.history.get('val_accuracy'):
    final_best_val_acc = max(history_fine_tune.history['val_accuracy'])
    print(f"Best validation accuracy during fine-tuning phase: {final_best_val_acc}")
else:
    print("Fine-tuning phase completed, but no validation accuracy history recorded")
















print(f"Model training complete. Final best model saved to {MODEL_SAVE_PATH}")

```

```

--- Starting Fine-Tuning Training for up to 15 epochs ---
Continuing training from epoch 5.
Saving best model checkpoints during this phase to: /cluster/home/miolate21/
FER_biasmitigation1/models/resnet50_baseline_224.keras
Training will run until epoch 20 unless stopped early.
Epoch 6/20
1037/1037  0s 67ms/step - accuracy: 0.1781 - loss: 2.722
2
Epoch 6: val_accuracy improved from -inf to 0.40248, saving model to /cluste
r/home/miolate21/FER_biasmitigation1/models/resnet50_baseline_224.keras
1037/1037  99s 87ms/step - accuracy: 0.1781 - loss: 2.72
19 - val_accuracy: 0.4025 - val_loss: 1.7137
Epoch 7/20
1036/1037  0s 63ms/step - accuracy: 0.3317 - loss: 1.981
4
Epoch 7: val_accuracy improved from 0.40248 to 0.46348, saving model to /clu
ster/home/miolate21/FER_biasmitigation1/models/resnet50_baseline_224.keras
1037/1037  81s 78ms/step - accuracy: 0.3317 - loss: 1.98
12 - val_accuracy: 0.4635 - val_loss: 1.4855
Epoch 8/20
1036/1037  0s 62ms/step - accuracy: 0.3908 - loss: 1.720
5
Epoch 8: val_accuracy improved from 0.46348 to 0.48855, saving model to /clu
ster/home/miolate21/FER_biasmitigation1/models/resnet50_baseline_224.keras
1037/1037  82s 79ms/step - accuracy: 0.3908 - loss: 1.72
04 - val_accuracy: 0.4885 - val_loss: 1.3934
Epoch 9/20
1036/1037  0s 63ms/step - accuracy: 0.4244 - loss: 1.599
7
Epoch 9: val_accuracy improved from 0.48855 to 0.49192, saving model to /clu
ster/home/miolate21/FER_biasmitigation1/models/resnet50_baseline_224.keras
1037/1037  82s 79ms/step - accuracy: 0.4244 - loss: 1.59
96 - val_accuracy: 0.4919 - val_loss: 1.3682
Epoch 10/20
1037/1037  0s 62ms/step - accuracy: 0.4556 - loss: 1.498
3
Epoch 10: val_accuracy improved from 0.49192 to 0.51242, saving model to /cl
uster/home/miolate21/FER_biasmitigation1/models/resnet50_baseline_224.keras
1037/1037  81s 78ms/step - accuracy: 0.4557 - loss: 1.49
83 - val_accuracy: 0.5124 - val_loss: 1.3167
Epoch 11/20
1036/1037  0s 63ms/step - accuracy: 0.4824 - loss: 1.413
4
Epoch 11: val_accuracy did not improve from 0.51242
1037/1037  80s 77ms/step - accuracy: 0.4824 - loss: 1.41
34 - val_accuracy: 0.4981 - val_loss: 1.3316
Epoch 12/20
1036/1037  0s 62ms/step - accuracy: 0.5006 - loss: 1.364
8
Epoch 12: val_accuracy improved from 0.51242 to 0.52290, saving model to /cl
uster/home/miolate21/FER_biasmitigation1/models/resnet50_baseline_224.keras
1037/1037  81s 78ms/step - accuracy: 0.5006 - loss: 1.36
48 - val_accuracy: 0.5229 - val_loss: 1.2773
Epoch 13/20
1037/1037  0s 63ms/step - accuracy: 0.5185 - loss: 1.323
2

```

Epoch 13: val_accuracy improved from 0.52290 to 0.53701, saving model to /cluster/home/miolate21/FER_biasmitigation1/models/resnet50_baseline_224.keras
1037/1037  **82s** 79ms/step - accuracy: 0.5185 - loss: 1.3232 - val_accuracy: 0.5370 - val_loss: 1.2337
Epoch 14/20
1036/1037  **0s** 63ms/step - accuracy: 0.5299 - loss: 1.2806
Epoch 14: val_accuracy improved from 0.53701 to 0.54508, saving model to /cluster/home/miolate21/FER_biasmitigation1/models/resnet50_baseline_224.keras
1037/1037  **82s** 79ms/step - accuracy: 0.5299 - loss: 1.2806 - val_accuracy: 0.5451 - val_loss: 1.2139
Epoch 15/20
1037/1037  **0s** 62ms/step - accuracy: 0.5333 - loss: 1.2623
Epoch 15: val_accuracy did not improve from 0.54508
1037/1037  **79s** 76ms/step - accuracy: 0.5333 - loss: 1.2623 - val_accuracy: 0.5371 - val_loss: 1.2246
Epoch 16/20
1036/1037  **0s** 63ms/step - accuracy: 0.5453 - loss: 1.2241
Epoch 16: val_accuracy did not improve from 0.54508
1037/1037  **80s** 77ms/step - accuracy: 0.5453 - loss: 1.2241 - val_accuracy: 0.5421 - val_loss: 1.2085
Epoch 17/20
1036/1037  **0s** 62ms/step - accuracy: 0.5606 - loss: 1.1843
Epoch 17: val_accuracy improved from 0.54508 to 0.54713, saving model to /cluster/home/miolate21/FER_biasmitigation1/models/resnet50_baseline_224.keras
1037/1037  **81s** 78ms/step - accuracy: 0.5606 - loss: 1.1843 - val_accuracy: 0.5471 - val_loss: 1.2082
Epoch 18/20
1037/1037  **0s** 62ms/step - accuracy: 0.5597 - loss: 1.1908
Epoch 18: val_accuracy improved from 0.54713 to 0.56714, saving model to /cluster/home/miolate21/FER_biasmitigation1/models/resnet50_baseline_224.keras
1037/1037  **81s** 78ms/step - accuracy: 0.5597 - loss: 1.1908 - val_accuracy: 0.5671 - val_loss: 1.1594
Epoch 19/20
1036/1037  **0s** 62ms/step - accuracy: 0.5793 - loss: 1.1497
Epoch 19: val_accuracy improved from 0.56714 to 0.57027, saving model to /cluster/home/miolate21/FER_biasmitigation1/models/resnet50_baseline_224.keras
1037/1037  **81s** 78ms/step - accuracy: 0.5793 - loss: 1.1497 - val_accuracy: 0.5703 - val_loss: 1.1583
Epoch 20/20
1036/1037  **0s** 62ms/step - accuracy: 0.5861 - loss: 1.1168
Epoch 20: val_accuracy improved from 0.57027 to 0.57317, saving model to /cluster/home/miolate21/FER_biasmitigation1/models/resnet50_baseline_224.keras
1037/1037  **81s** 78ms/step - accuracy: 0.5860 - loss: 1.1168 - val_accuracy: 0.5732 - val_loss: 1.1464
Restoring model weights from the end of the best epoch: 20.

--- Fine-tuning training phase finished. ---

Best validation accuracy during fine-tuning phase: 0.5732

Model training complete. Final best model saved to /cluster/home/miolate21/F

ER_biasmitigation1/models/resnet50_baseline_224.keras (based on validation accuracy).

```
In [53]: # Cell 10b Continue Fine-Tuning Training

# --- Configuration for Continued Training ---
CONTINUE_FROM_EPOCH = 20
ADDITIONAL_EPOCHS = 10
NEW_TOTAL_EPOCHS = CONTINUE_FROM_EPOCH + ADDITIONAL_EPOCHS

MODEL_SAVE_PATH = "/cluster/home/miolate21/FER_biasmitigation1/models/resnet50_baseline_224.keras"
FINE_TUNE_LR = 1e-5

print(f"\n--- Continuing Fine-Tuning Training for {ADDITIONAL_EPOCHS} more epochs")
print(f"Starting from epoch {CONTINUE_FROM_EPOCH + 1} (running up to epoch {NEW_TOTAL_EPOCHS})")
print(f"Saving best model checkpoints during this phase to: {MODEL_SAVE_PATH}")

print("Defining fresh callbacks for continued training.")
callbacks_continue_finetune = [
    tf.keras.callbacks.ModelCheckpoint(
        filepath=MODEL_SAVE_PATH,
        save_best_only=True,
        monitor='val_accuracy',
        mode='max',
        verbose=1
    ),
    tf.keras.callbacks.EarlyStopping(
        monitor='val_accuracy',
        patience=5,
        verbose=1,
        restore_best_weights=True
    )
]

if 'model' not in globals():
    print("WARN: 'model' object not found. Attempting to load from MODEL_SAVE_PATH")
    try:
        model = tf.keras.models.load_model(MODEL_SAVE_PATH)
        print("Model loaded successfully.")
    except Exception as e:
        raise NameError(f"ERROR: 'model' object not found and failed to load from {MODEL_SAVE_PATH}")

if 'train_ds' not in globals() or 'val_ds' not in globals() or train_ds is None or val_ds is None:
    raise NameError("ERROR: 'train_ds' or 'val_ds' not defined or is None.")

if tf.data.experimental.cardinality(train_ds).numpy() == 0 or \
    tf.data.experimental.cardinality(val_ds).numpy() == 0:
    raise ValueError("ERROR: Training or validation dataset is empty.")

print(f"Model is currently compiled with optimizer: {model.optimizer.name}")
if abs(model.optimizer.learning_rate.numpy() - FINE_TUNE_LR) > 1e-9:
    print(f"WARN: Model LR ({model.optimizer.learning_rate.numpy():.1E}) differs from FINE_TUNE_LR ({FINE_TUNE_LR:.1E})")

# --- Continue Training ---
print(f"\nCalling model.fit from epoch {CONTINUE_FROM_EPOCH} up to {NEW_TOTAL_EPOCHS}")
history_continue_fine_tune = model.fit(
    train_ds,
    validation_data=val_ds,
```



```

epochs=NEW_TOTAL_EPOCHS,
initial_epoch=CONTINUE_FROM_EPOCH,
verbose=1
)

print(f"Model training complete (extended). Final best model saved to {MODEL

```

--- Continuing Fine-Tuning Training for 10 more epochs ---

Starting from epoch 21 (running up to epoch 30).


Saving best model checkpoints during this phase to: /cluster/home/miolate21/ FER_biasmitigation1/models/resnet50_baseline_224.keras

Defining fresh callbacks for continued training.


Model is currently compiled with optimizer: adam, LR: 1.0E-05

Calling model.fit from epoch 20 up to 30...


Epoch 21/30

1037/1037  **78s** 75ms/step - accuracy: 0.5921 - loss: 1.1146 - val_accuracy: 0.5769 - val_loss: 1.1337


Epoch 22/30

1037/1037  **78s** 76ms/step - accuracy: 0.5990 - loss: 1.0912 - val_accuracy: 0.5800 - val_loss: 1.1211


Epoch 23/30

1037/1037  **79s** 76ms/step - accuracy: 0.6061 - loss: 1.0771 - val_accuracy: 0.5851 - val_loss: 1.1217


Epoch 24/30

1037/1037  **78s** 76ms/step - accuracy: 0.6092 - loss: 1.0556 - val_accuracy: 0.5875 - val_loss: 1.1233


Epoch 25/30

1037/1037  **79s** 76ms/step - accuracy: 0.6154 - loss: 1.0438 - val_accuracy: 0.5881 - val_loss: 1.1082


Epoch 26/30

1037/1037  **78s** 76ms/step - accuracy: 0.6235 - loss: 1.0205 - val_accuracy: 0.5947 - val_loss: 1.1026


Epoch 27/30

1037/1037  **79s** 76ms/step - accuracy: 0.6272 - loss: 1.0030 - val_accuracy: 0.5963 - val_loss: 1.0941


Epoch 28/30

1037/1037  **78s** 75ms/step - accuracy: 0.6336 - loss: 1.0011 - val_accuracy: 0.5990 - val_loss: 1.0861

Epoch 29/30

1037/1037  **78s** 75ms/step - accuracy: 0.6410 - loss: 0.9778 - val_accuracy: 0.5956 - val_loss: 1.0900

Epoch 30/30

1037/1037  **79s** 76ms/step - accuracy: 0.6441 - loss: 0.9639 - val_accuracy: 0.5956 - val_loss: 1.0972

Model training complete (extended). Final best model saved to /cluster/home/miolate21/ FER_biasmitigation1/models/resnet50_baseline_224.keras.

In [56]: *#Cell 11 Plotting Combined Training History*

```

import matplotlib.pyplot as plt

print("\n--- Plotting Combined Training History ---")

acc = []
val_acc = []

```

```

loss = []
val_loss = []
total_epochs_run = 0

if 'history_initial' in globals() and hasattr(history_initial, 'history'):
    acc += history_initial.history.get('accuracy', [])
    val_acc += history_initial.history.get('val_accuracy', [])
    loss += history_initial.history.get('loss', [])
    val_loss += history_initial.history.get('val_loss', [])
    total_epochs_run += len(history_initial.history.get('accuracy', []))
    print(f"Initial history loaded: {len(history_initial.history.get('accuracy', []))}")

if 'history_fine_tune' in globals() and hasattr(history_fine_tune, 'history'):
    acc += history_fine_tune.history.get('accuracy', [])
    val_acc += history_fine_tune.history.get('val_accuracy', [])
    loss += history_fine_tune.history.get('loss', [])
    val_loss += history_fine_tune.history.get('val_loss', [])
    total_epochs_run += len(history_fine_tune.history.get('accuracy', []))
    print(f"First fine-tune history loaded: {len(history_fine_tune.history.get('accuracy', []))}")

if 'history_continue_fine_tune' in globals() and hasattr(history_continue_fine_tune, 'history'):
    acc += history_continue_fine_tune.history.get('accuracy', [])
    val_acc += history_continue_fine_tune.history.get('val_accuracy', [])
    loss += history_continue_fine_tune.history.get('loss', [])
    val_loss += history_continue_fine_tune.history.get('val_loss', [])
    total_epochs_run += len(history_continue_fine_tune.history.get('accuracy', []))
    print(f"Continued fine-tune history loaded: {len(history_continue_fine_tune.history.get('accuracy', []))}")

if not acc:
    print("ERROR: No training history data found. Cannot plot.")
else:
    epochs_range = range(total_epochs_run)

    plt.figure(figsize=(14, 6))

    plt.subplot(1, 2, 1)
    plt.plot(epochs_range, acc, label='Training Accuracy', alpha=0.8)
    plt.plot(epochs_range, val_acc, label='Validation Accuracy', linewidth=1)
    initial_epochs_count = len(history_initial.history.get('accuracy', []))
    first_finetune_epochs_count = len(history_fine_tune.history.get('accuracy', []))

    if initial_epochs_count > 0:
        plt.axvline(initial_epochs_count - 1, linestyle='--', color='r', label='Initial Epochs')
    if first_finetune_epochs_count > 0 and initial_epochs_count > 0:
        plt.axvline(initial_epochs_count + first_finetune_epochs_count - 1, linestyle='--', color='b', label='First Finetune Epochs')

    plt.legend(loc='lower right')
    plt.title('Training and Validation Accuracy')
    plt.xlabel('Epoch')
    plt.ylabel('Accuracy')
    plt.grid(True, linestyle='--', alpha=0.6)

    plt.subplot(1, 2, 2)
    plt.plot(epochs_range, loss, label='Training Loss', alpha=0.8)

```

```

plt.plot(epochs_range, val_loss, label='Validation Loss', linewidth=1.5)
if initial_epochs_count > 0:
    plt.axvline(initial_epochs_count - 1, linestyle='--', color='r', label='Initial Epochs')
if first_finetune_epochs_count > 0 and initial_epochs_count > 0:
    plt.axvline(initial_epochs_count + first_finetune_epochs_count - 1,
                 linestyle='--', color='g', label='Finetune Epochs')

plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.grid(True, linestyle='--', alpha=0.6)

plt.tight_layout()

directory = '/cluster/home/miolate21/FER_biasmitigation1/results'
filename = 'resnet50_training_metrics_e30.png'
filepath = os.path.join(directory, filename)

plt.savefig(filepath)

plt.show()

print("\n--- Evaluating Final Model ---")

if 'val_ds' in globals() and val_ds is not None:
    print("Evaluating model on validation set...")
    eval_loss, eval_acc = model.evaluate(val_ds, verbose=1)
    print(f"\nFinal evaluation on validation set:")
    print(f"Validation Loss: {eval_loss:.4f}")
    print(f"Validation Accuracy: {eval_acc:.4f}")

    print("\nGenerating classification report and confusion matrix...")
    import numpy as np
    from sklearn.metrics import classification_report, confusion_matrix
    import seaborn as sns

    y_pred_probs = model.predict(val_ds)
    y_pred = np.argmax(y_pred_probs, axis=1)
    y_true = np.concatenate([y for x, y in val_ds], axis=0)

    if 'CLASS_NAMES' not in globals() or CLASS_NAMES is None:
        print("Warning: CLASS_NAMES not found, using default labels.")
        target_names = [str(i) for i in range(NUM_CLASSES)]
    else:
        target_names = CLASS_NAMES

    print("\nClassification Report:")
    print(classification_report(y_true, y_pred, target_names=target_names))

    cm = confusion_matrix(y_true, y_pred)
    plt.figure(figsize=(9, 7))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=target_names,
                yticklabels=target_names)
    plt.xlabel('Predicted Label')
    plt.ylabel('True Label')

```

```

plt.title('Confusion Matrix')

directory = '/cluster/home/miolate21/FER_biasmitigation1/results'
filename = 'resnet50_confusion_matrix.png'
filepath = os.path.join(directory, filename)

plt.savefig(filepath)

plt.show()
else:
    print("Validation dataset ('val_ds') not found. Skipping final evaluation")


```

--- Plotting Combined Training History ---

Initial history loaded: 5 epochs.


First fine-tune history loaded: 15 epochs.

Continued fine-tune history loaded: 10 epochs.



--- Evaluating Final Model ---

Evaluating model on validation set...


260/260  **14s** 54ms/step - accuracy: 0.6007 - loss: 1.0730

Final evaluation on validation set:

Validation Loss: 1.0972

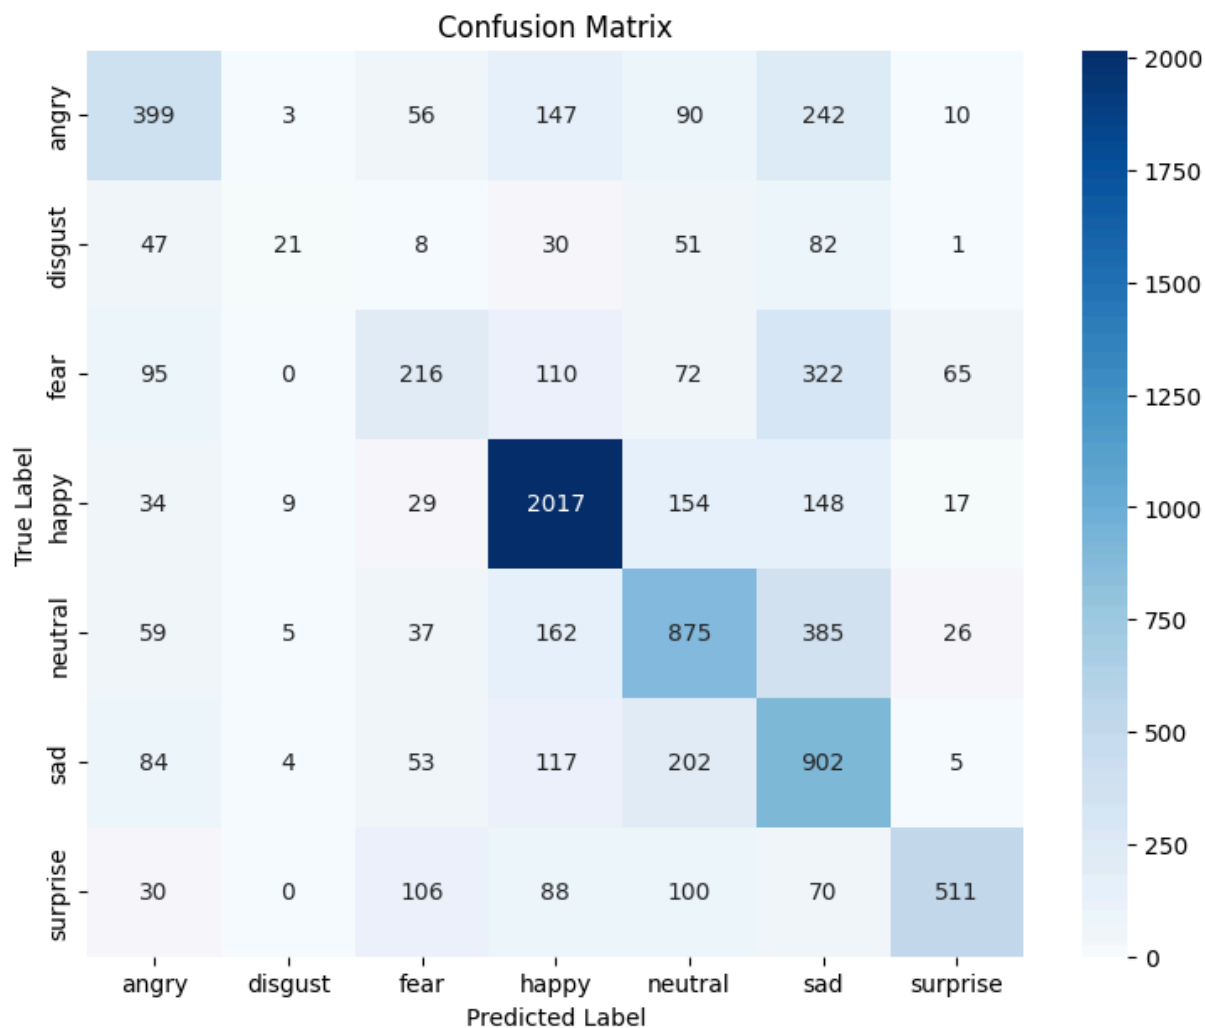
Validation Accuracy: 0.5956

Generating classification report and confusion matrix...

260/260  **15s** 56ms/step

Classification Report:

	precision	recall	f1-score	support
angry	0.53	0.42	0.47	947
disgust	0.50	0.09	0.15	240
fear	0.43	0.25	0.31	880
happy	0.76	0.84	0.79	2408
neutral	0.57	0.56	0.57	1549
sad	0.42	0.66	0.51	1367
surprise	0.80	0.56	0.66	905
accuracy			0.60	8296
macro avg	0.57	0.48	0.50	8296
weighted avg	0.60	0.60	0.58	8296



```
In [57]: from sklearn.metrics import classification_report
import os

# Generate classification report
report_text = classification_report(y_true, y_pred, target_names=CLASS_NAMES)

# === Dynamic file naming ===
model_version = "resnet50_tl_finetune"
total_epochs = len(acc)
filename = f"classification_report_baseline_{model_version}_e{total_epochs}."

# === Save path ===
report_dir = '/cluster/home/miolate21/FER_biasmitigation1/results'
report_path = os.path.join(report_dir, filename)

# Save the report
with open(report_path, 'w') as f:
    f.write(f"Classification Report - Baseline Model ({model_version})\n\n")
    f.write(report_text)

print(f"Classification report saved to: {report_path}")
```

Classification report saved to: /cluster/home/miolate21/FER_biasmitigation1/results/classification_report_baseline_resnet50_tl_finetune_e30.txt

In []: